

availability

UFRN



Agenda

- 1 10:00 AM: Session I - Introduction to PaScal Suite
 - Performance versus Scalability
 - Introduction to PaScal Analyzer
 - Introduction to PaScal Viewer
- 2 11:00 AM: Coffee break
- 3 11:30 AM: Session II - Hands-on (bring your laptop computer)
 - Evaluating scalability with PaScal Suite
 - Profiling an OpenMP binary for scalability
 - Instrumenting source code manually
 - Final remarks



Agenda

- 1 10:00 AM: Session I - Introduction to PaScal Suite
 - Performance versus Scalability
 - Introduction to PaScal Analyzer
 - Introduction to PaScal Viewer
- 2 11:00 AM: Coffee break
- 3 11:30 AM: Session II - Hands-on (bring your laptop computer)
 - Evaluating scalability with PaScal Suite
 - Profiling an OpenMP binary for scalability
 - Instrumenting source code manually
 - Final remarks

Performance versus Scalability

Objectives of this tutorial

- ▶ Present the difference between performance and scalability
- ▶ Show existing performance-profile tools are unsuited for profiling scalability
- ▶ Emphasize the growing need for scalable parallel programs
- ▶ Demonstrate how PaScal Suite can assist developers with scalability analysis
- ▶ Teach in practice how to use it to identify scalability hotspots and bottlenecks

Performance versus Scalability

A metaphor for a program



- ▶ Performance is related to how fast the task is executed
- ▶ Scalability is concerned with how efficiently the task is performed

Performance versus Scalability

Here:

- ▶ Performance is the inverse of Time, i.e.
 - ▶ $\text{Performance} = 1/\text{Time}$
- ▶ Efficiency is how much work is performed per resource, i.e.
 - ▶ $\text{Efficiency} = \text{Work} \times \text{Performance} / \text{Resources}$, or
 - ▶ $\text{Efficiency} = \text{Work} / (\text{Time} \times \text{Resources})$

The unity task performed with one resource in the unity time is 100% efficient



Work = 1
Time = 1
Resource = 1
Efficiency = 1

Performance versus Scalability

The first task is sequential, and the second is parallel. Both perform 100% efficiently.



Work = 8
Time = 8
Resource = 1
Efficiency = 1



Work = 8
Time = 4
Resource = 2
Efficiency = $W/(T \times R) = 1$




Note: Resources can scale for the sequential task as well, which would relate to using more sophisticated hardware. Mixing both types of resource scaling would relate to heterogeneous processing.

Performance versus Scalability

Scaling up resources alone often harms efficiency



Work = 8
Time = 3
Resource = 3
Efficiency = $8/(3 \times 3) = 0.89$

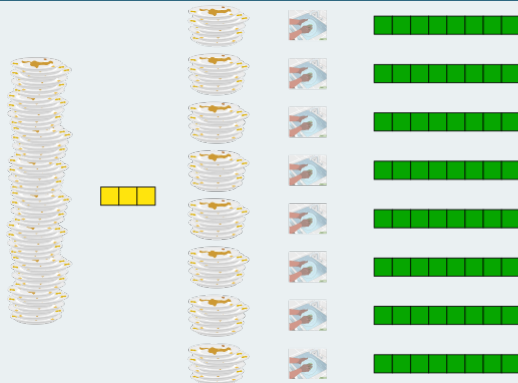


Work = 8
Time = 1
Resource = 20
Efficiency = $8/(1 \times 20) = 0.4$

Performance versus Scalability

Larger machines, like supercomputers, require larger problems to sustain efficiency

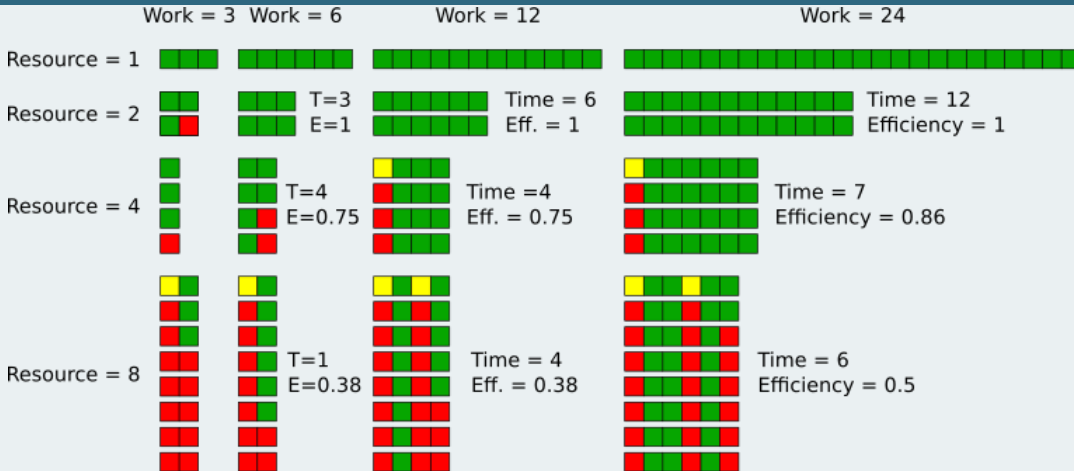
Parallelization overhead, dependencies, serial parts, and so on also harms efficiency



Work = 64
Time = 3 + 8
Resource = 8
Efficiency =
 $64 / (11 * 8) = 0.73$

Performance versus Scalability

Efficiency can vary with resource scaling and problem scaling



Performance versus Scalability

- ▶ Scalability analysis requires coarse evaluation of many scaling configurations
- ▶ Performance analysis involves a detailed evaluation of a given configuration

How unsuited are existing performance-profiling tools to assess scalability


- ▶ Require the user to run and collect information on many configurations manually
- ▶ Provide much detailed data beyond the necessary to evaluate scalability
- ▶ Have no native visual artifacts that guide and facilitate the scalability analysis

Despite the shortage of suitable analysis tools, parallel scalability is a growing concern.




Performance versus Scalability

- ▶ Top 500 difference between peak, HPL, and HPGC performance.
- ▶ Very large machines are good for solving many problems at the same time but are not efficient to solve very large problems.



Single node scaling soon to face harder scaling problems


- ▶ Moore's law lives
- ▶ Single thread big slow down
- ▶ MHz plateaued
- ▶ Number of cores follows Moore's law



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
Plot and data collected for 2010–2019 by K. Rupp


Sources of inefficiency in (parallel) computing

- ▶ Due to:
 - ▶ Load imbalance
 - ▶ Synchronization
 - ▶ Communication
- ▶ Lead to:
 - ▶ Extra computation
 - ▶ Idle time
- ▶ Optimization approaches
 - ▶ avoid costs
- ▶ Trade-off approaches
 - ▶ Trade costs




Sources of inefficiency in (parallel) computing

- ▶ Due to:
 - ▶ Load imbalance
 - ▶ Synchronization
 - ▶ Communication
- ▶ Lead to:
 - ▶ Extra computation
 - ▶ Idle time
- ▶ Optimization approaches
 - ▶ avoid costs
- ▶ Trade-off approaches
 - ▶ Trade costs




Sources of inefficiency in (parallel) computing

- ▶ Due to:
 - ▶ Load imbalance
 - ▶ Synchronization
 - ▶ Communication
- ▶ Lead to:
 - ▶ Extra computation
 - ▶ Idle time
- ▶ Optimization approaches
 - ▶ avoid costs
- ▶ Trade-off approaches
 - ▶ Trade costs




Speedup models

Amdahl's model




Speedup models

Amdahl's model

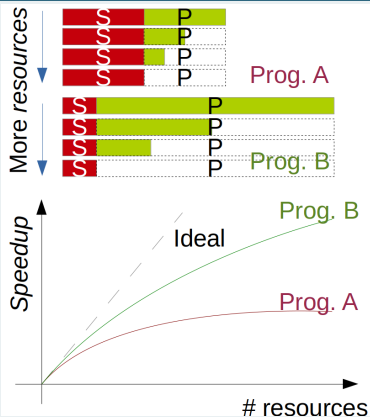


Gustafson's model



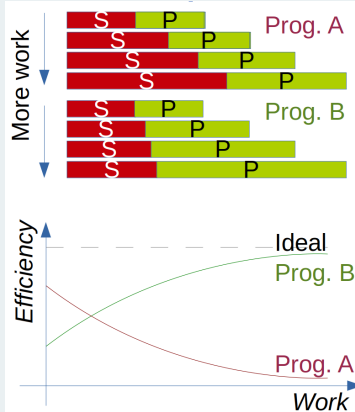
Speedup models

Amdahl's model



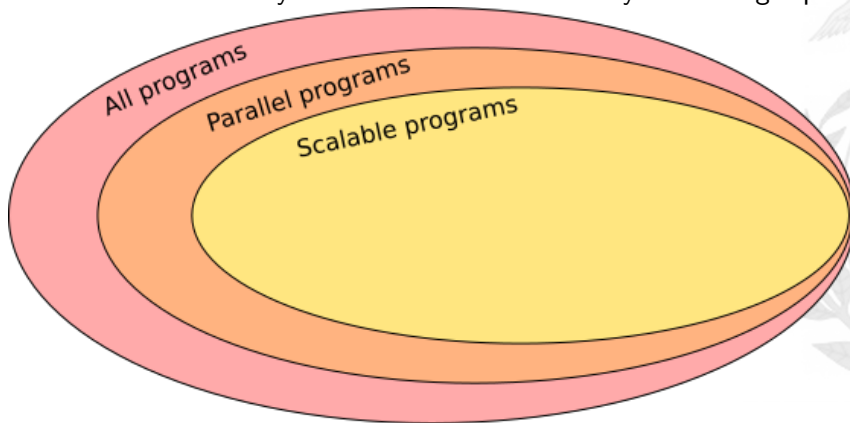
► Many other models

Gustafson's model



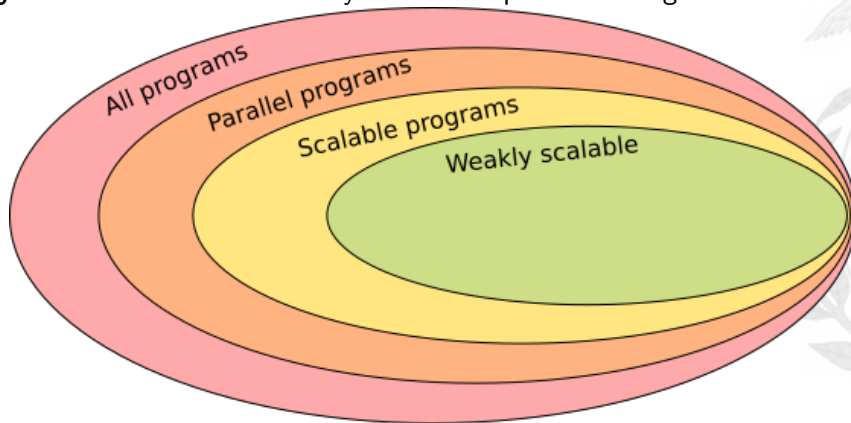
Parallel Efficiency and Scalability Analysis

Scalable: sustain efficiency for more resources with any-factor larger problem



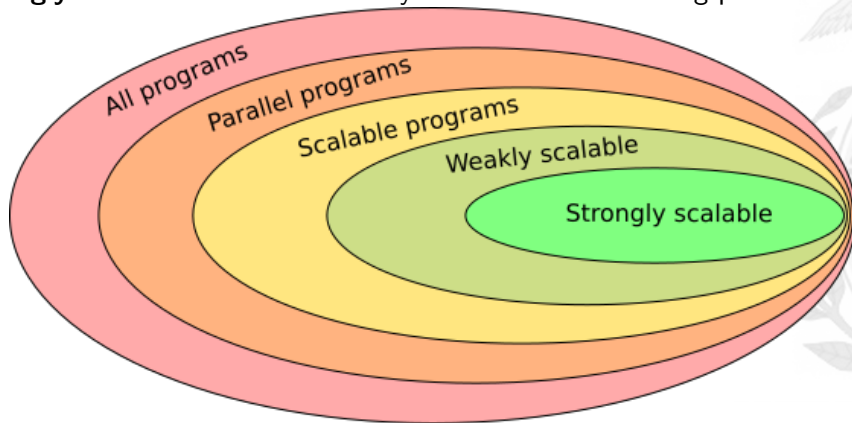
Parallel Efficiency and Scalability Analysis

Weakly scalable: Sustain efficiency even when problem size grows less than resources



Parallel Efficiency and Scalability Analysis

Strongly scalable: Sustain efficiency even without increasing problem size



Assessing a program's scalability with tables

$$T_1 = n^2$$

$$T_P = \frac{n^2}{p} + \log_2 p$$

Timings

| # cores | Problem size | | | | | | |
|---------|--------------|--------|--------|---------|---------|----------|-----------|
| | 1x | 2x | 4x | 8X | 16X | 32X | |
| | 1 | 100,00 | 400,00 | 1600,00 | 6400,00 | 25600,00 | 102400,00 |
| | 2 | 51,00 | 201,00 | 801,00 | 3201,00 | 12801,00 | 51201,00 |
| | 4 | 27,00 | 102,00 | 402,00 | 1602,00 | 6402,00 | 25602,00 |
| | 8 | 15,50 | 53,00 | 203,00 | 803,00 | 3203,00 | 12803,00 |
| | 16 | 10,25 | 29,00 | 104,00 | 404,00 | 1604,00 | 6404,00 |
| | 32 | 8,13 | 17,50 | 55,00 | 205,00 | 805,00 | 3205,00 |
| | 64 | 7,56 | 12,25 | 31,00 | 106,00 | 406,00 | 1606,00 |
| 128 | 7,78 | 10,13 | 19,50 | 57,00 | 207,00 | 807,00 | |

Assessing a program's scalability with tables

$$T_1 = n^2$$

$$T_P = \frac{n^2}{p} + \log_2 p$$

$$= \frac{T_1}{T_p}$$

Speedups

| | Problem size | | | | | |
|---------|--------------|-------|-------|-------|--------|--------|
| | 1x | 2x | 4x | 8X | 16X | 32X |
| # cores | 1 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 |
| | 2 | 1,96 | 1,99 | 2,00 | 2,00 | 2,00 |
| | 4 | 3,70 | 3,92 | 3,98 | 4,00 | 4,00 |
| | 8 | 6,45 | 7,55 | 7,88 | 7,97 | 8,00 |
| | 16 | 9,76 | 13,79 | 15,38 | 15,84 | 15,99 |
| | 32 | 12,31 | 22,86 | 29,09 | 31,22 | 31,95 |
| | 64 | 13,22 | 32,65 | 51,61 | 60,38 | 63,76 |
| | 128 | 12,85 | 39,51 | 82,05 | 112,28 | 126,89 |

Assessing a program's scalability with tables

$$T_1 = n^2$$

$$T_P = \frac{n^2}{p} + \log_2 p$$

$$= \frac{T_1}{T_p}$$

$$E = \frac{S}{p}$$

Efficiencies

| | Problem size | | | | | |
|------------|--------------|------|------|------|------|------|
| | 1x | 2x | 4x | 8X | 16X | 32X |
| 1 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 |
| 2 | 0,98 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 |
| 4 | 0,93 | 0,98 | 1,00 | 1,00 | 1,00 | 1,00 |
| 8 | 0,81 | 0,94 | 0,99 | 1,00 | 1,00 | 1,00 |
| 16 | 0,61 | 0,86 | 0,96 | 0,99 | 1,00 | 1,00 |
| 32 | 0,38 | 0,71 | 0,91 | 0,98 | 0,99 | 1,00 |
| 64 | 0,21 | 0,51 | 0,81 | 0,94 | 0,99 | 1,00 |
| 128 | 0,10 | 0,31 | 0,64 | 0,88 | 0,97 | 0,99 |

cores

Assessing a program's scalability with tables

$$T_1 = n^2$$

$$T_P = \frac{n^2}{p} + \log_2 p$$

$$= \frac{T_1}{T_p}$$

$$E = \frac{S}{p}$$

Efficiencies

| | Problem size | | | | | |
|---------|--------------|------|------|------|------|------|
| | 1x | 2x | 4x | 8X | 16X | 32X |
| # cores | 1 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 |
| | 2 | 0,98 | 1,00 | 1,00 | 1,00 | 1,00 |
| | 4 | 0,93 | 0,98 | 1,00 | 1,00 | 1,00 |
| | 8 | 0,81 | 0,94 | 0,99 | 1,00 | 1,00 |
| | 16 | 0,61 | 0,86 | 0,96 | 0,99 | 1,00 |
| | 32 | 0,38 | 0,71 | 0,91 | 0,98 | 1,00 |
| | 64 | 0,21 | 0,51 | 0,81 | 0,94 | 1,00 |
| | 128 | 0,10 | 0,31 | 0,64 | 0,88 | 0,99 |

Scalable? Weakly scalable? Or strongly scalable?

Assessing a program's scalability with tables

$$T_1 = n^2$$

$$T_P = \frac{n^2}{p} + \log_2 p$$

$$= \frac{T_1}{T_p}$$

$$E = \frac{S}{p}$$

Efficiencies

| | Problem size | | | | | |
|---------|--------------|------|------|------|------|------|
| | 1x | 2x | 4x | 8X | 16X | 32X |
| # cores | 1 | 1,00 | 1,00 | 1,00 | 1,00 | 1,00 |
| | 2 | 0,98 | 1,00 | 1,00 | 1,00 | 1,00 |
| | 4 | 0,93 | 0,98 | 1,00 | 1,00 | 1,00 |
| | 8 | 0,81 | 0,94 | 0,99 | 1,00 | 1,00 |
| | 16 | 0,61 | 0,86 | 0,96 | 0,99 | 1,00 |
| | 32 | 0,38 | 0,71 | 0,91 | 0,98 | 1,00 |
| | 64 | 0,21 | 0,51 | 0,81 | 0,94 | 1,00 |
| | 128 | 0,10 | 0,31 | 0,64 | 0,88 | 0,99 |

Scalable? Weakly scalable? Or strongly scalable?

Analysis still doable with small core count and quadratically increasing rates

Assessing a program's scalability with tables

| | Problem size | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---------|--------------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|
| | 1x | 2x | 3x | 4x | 5x | 6x | 7x | 8x | 9x | 10x | 11x | 12x | 13x | 14x | 15x | 16x | 17x | 18x | 19x | 20x | 21x | 22x | 23x | 24x | 25x | 26x | 27x | 28x | 29x |
| # cores | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 |
| 1 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 2 | 0.98 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 3 | 0.95 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 4 | 0.93 | 0.98 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 5 | 0.90 | 0.97 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 6 | 0.87 | 0.96 | 0.98 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 7 | 0.84 | 0.95 | 0.98 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 8 | 0.81 | 0.94 | 0.97 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 9 | 0.78 | 0.93 | 0.97 | 0.98 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 10 | 0.75 | 0.92 | 0.96 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 11 | 0.72 | 0.91 | 0.96 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 12 | 0.70 | 0.90 | 0.95 | 0.97 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 13 | 0.68 | 0.89 | 0.95 | 0.97 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 14 | 0.65 | 0.88 | 0.94 | 0.97 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 15 | 0.63 | 0.87 | 0.94 | 0.96 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 16 | 0.61 | 0.86 | 0.93 | 0.96 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 17 | 0.59 | 0.85 | 0.93 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 18 | 0.57 | 0.84 | 0.92 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 19 | 0.55 | 0.83 | 0.92 | 0.95 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 20 | 0.54 | 0.82 | 0.91 | 0.95 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 21 | 0.52 | 0.81 | 0.91 | 0.95 | 0.96 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 22 | 0.50 | 0.80 | 0.90 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 23 | 0.49 | 0.79 | 0.90 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 24 | 0.48 | 0.78 | 0.89 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 25 | 0.46 | 0.78 | 0.89 | 0.93 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 26 | 0.45 | 0.77 | 0.88 | 0.93 | 0.95 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 27 | 0.44 | 0.76 | 0.88 | 0.93 | 0.95 | 0.97 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 28 | 0.43 | 0.75 | 0.87 | 0.92 | 0.95 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 29 | 0.42 | 0.74 | 0.86 | 0.92 | 0.95 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 30 | 0.40 | 0.73 | 0.86 | 0.92 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 31 | 0.39 | 0.72 | 0.85 | 0.91 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |
| 32 | 0.38 | 0.71 | 0.85 | 0.91 | 0.94 | 0.96 | 0.97 | 0.98 | 0.98 | 0.98 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 0.99 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 | 1.00 |

1/4 of the table of the previous slide with linear rates

Assessing a program's scalability with tables

- ▶ Building tables manually is tedious and time-consuming
 - ▶ Remember: performance-profiling tools do not build these tables. They generate data (way too much) for a single table cell.
 - ▶ The scalability-caring developer is left alone to run, collect and, somehow, find a way to visualize (scalability) hot spots and bottlenecks
- ▶ Tables hide scalability trends
- ▶ Line or bar plots are unsuited for visualizing this type of data. Cloaks the trends.
- ▶ What about profiling for scalability?
 - ▶ Requires a table for each code region of interest
 - ▶ Manually: more tedious and time-consuming
 - ▶ Using profiling tools: even more tedious and time-consuming

Introduction to PaScal Analyzer

PaScal Analyzer

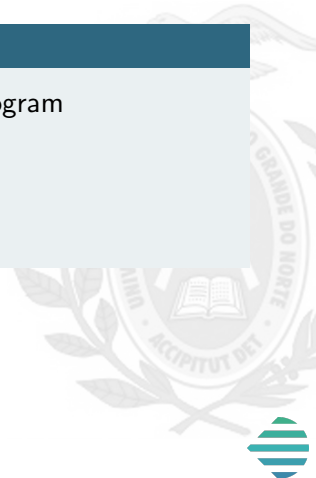
A tool and library from the Parallel Scalability Suite toolset to control, measure, and collect information from multiple executions of a parallel application automatically.

Note: Currently suited for shared-memory environments.

Introduction to PaScal Analyzer

Objectives of the Parallel Scalability Analyzer

- ▶ Read the definitions of multiple configurations of a parallel program
- ▶ Measure data from those configurations automatically
- ▶ Collect the measured data in a non-intrusive way
- ▶ Organize the collected data for future visual analysis



Introduction to PaScal Analyzer

Objectives of the Parallel Scalability Analyzer

- ▶ Read the definitions of multiple configurations of a parallel program
- ▶ Measure data from those configurations automatically
- ▶ Collect the measured data in a non-intrusive way
- ▶ Organize the collected data for future visual analysis

Which data?

- ▶ The PaScal Analyzer can deal with timing, power, and performance counters
- ▶ In this tutorial, we will be only concerned with timing.

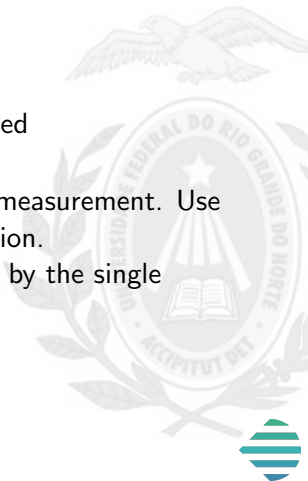
Requirements to use PaScal Analyzer

- ▶ Linux operating system
- ▶ For profiling:
 - ▶ applications written in C and C++
 - ▶ OpenMP or PThreads (no MPI yet)
 - ▶ profiling binary: GNU C Compiler



Instrumenting source code

- ▶ PThreads code and OpenMP code can be manually instrumented
 - ▶ Include the `pascalops.h` header file in our source files.
 - ▶ Use `pascal_start` to mark the beginning of a region for measurement. Use the single integer argument as the identification of the region.
 - ▶ Use `pascal_stop` to mark the end of the region identified by the single integer argument.



Instrumenting source code

Example of instrumentation of a OpenMP C program

```
#include "pascalops.h"
...
int main()
{
    ...
    pascal_start(1);
    #pragma omp parallel
    {
        ...
    }
    pascal_stop(1);
    ...
}
```

Compiling our code with the PaScalOps library

- ▶ After installation, the pascalops library should be visible to GCC
- ▶ Simply add the flag `-lmpascalops` for compiling manually instrumented sources.

Example compilation

```
$ gcc myapp.c -fopenmp -lmpascalops
```

Using PaScal Analyzer

Usage

```
pascalanalyzer [-h] [-c CORES] [-f FREQUENCIES] [-i INPUTS] [-o OUTPUT] [-r RPTS] [-g]
               [-t {aut,AUT,man,MAN}] [-a LEVEL] [--mpi RUNTIME] [--ragt TYPE] [-v VERBOSE]
               [--dcrs] [--dhpt] [--domp] [--dout] [--govr GOVERNOR] [--idtm TIME]
               [--fgpe EVENT] [--fgps EVENT] [--rple {sysfs,perf}]
               [--rpls {sysfs,scontrol,perf}] [--ipmi SERVER USER PASSWORD] [--modl NPTS MODE]
               [--prcs] [--lpcs] [--imnt PATH EXTENSIONS]
               [application]
```


Examples

```
$ pascalanalyzer ./myapp --inst aut --idtm 5 --cors 1:4 --ipts 10,20,30,40 --verb INFO
$ pascalanalyzer ./myapp -t man -g -r 5 -c 1,32 -i 10,320 -v 2 -o myoutput.json
```

Introduction to PaScal Analyzer

PaScal Viewer

A web tool from the Parallel Scalability Suite that loads PaScal Analyzer's outputs and presents four diagrams to assist the user in understanding the program's scaling behavior with visual elements that emphasize scalability hotspots and bottlenecks.





Agenda

- 1 10:00 AM: Session I - Introduction to PaScal Suite
 - Performance versus Scalability
 - Introduction to PaScal Analyzer
 - Introduction to PaScal Viewer
- 2 11:00 AM: Coffee break
- 3 11:30 AM: Session II - Hands-on (bring your laptop computer)
 - Evaluating scalability with PaScal Suite
 - Profiling an OpenMP binary for scalability
 - Instrumenting source code manually
 - Final remarks



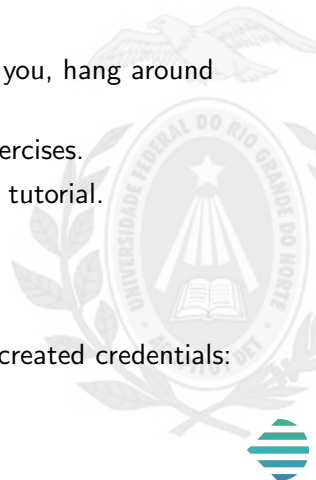
Agenda

- 1 10:00 AM: Session I - Introduction to PaScal Suite
 - Performance versus Scalability
 - Introduction to PaScal Analyzer
 - Introduction to PaScal Viewer
- 2 11:00 AM: Coffee break
- 3 11:30 AM: Session II - Hands-on (bring your laptop computer)
 - Evaluating scalability with PaScal Suite
 - Profiling an OpenMP binary for scalability
 - Instrumenting source code manually
 - Final remarks





Setting up your laptop

- ▶ This is a hands-on session, but if you do not have laptop with you, hang around to see PaScal Suite in action.
- ▶ We will use NPAD, UFRN's HPC cluster, to do the tutorial exercises.
- ▶ NPAD's team has reserved 3 nodes with 32 cores each for this tutorial.
- ▶ Access the link:
<http://npad.ufrn.br/npad/Tutorial>
and follow the steps to create your tutorial account.
- ▶ In order to login, use the following command with your newly created credentials:
`ssh -p4422 tutorial[YOUR NUMBER]@sc2.npad.ufrn.br`



Compile and evaluate the scalability of our application

- ing LU
- ke:
- for scaling the
- ur problem size,
- 
- 

Compile and evaluate the scalability of our application

- Use the following command to run your first analysis job:

```
sbatch pascalanalyzer_job.sh
```

- ▶ Inspect the status of your job with:

```
queue -u $USER
```

- Now, copy the output file from NPAD to your machine

```
scp -P 4422 tutorial[YOUR NUMBER]@sc2.npad.ufrn.br\
:pascal-suite-tutorial/*.json .
```

- And load it in the PaScaL Viewer webtool at the link:

<https://pascalsuite.imd.ufrn.br>

Profiling an OpenMP binary for scalability

- ▶ The previous analysis involved the whole program and the parallel regions of the code because the Analyzer identifies the OpenMP parallel regions in the binary as inner regions.
- ▶ To analyze the gaps between regions and generate a complete scalability profile, use the `-g` option.
- ▶ The `-g` option considers these parts of the program as regions of interest:
 - ▶ the region between the start of the program and the first identified region;
 - ▶ the gaps between identified regions; and
 - ▶ the region between the last identified region and the end of the program.
- ▶ Add the `-g` option in the `pascalanalyzer_job.sh` script, run the job again, and load the resulting `.json` file on the webtool again.

Instrumenting source code manually

- ▶ Include the header `pascalops.h` in the source `LU_rand_omp.cpp`
- ▶ Instrument the region(s) of your interest with:
 - ▶ `pascal_start(id)`, and
 - ▶ `pascal_stop(id)`,
 - ▶ where `id` is an integer that uniquely identifies a region.
- ▶ Then, include the flag `-lmpascalops` in the compilation line in the `Makefile` file to link our program with the PaScal library.
- ▶ Change the `-t` option in the job script from `aut` to `man`.
- ▶ Compile, submit the job, and visualize the result with PaScal Viewer again.



Final remarks

- ▶ **Disclaimer:** PaScal Suite is a prototype and should be used with no guarantees or liability to the authors.
- ▶ You can contribute by reporting bugs but we are also seeking partners to make further developments.
- ▶ Wish list of features:
 - ▶ MPI support
 - ▶ Reduce analysis costs by building scalability models from samples of the scaling space
 - ▶ Reduce analysis cost by using performance models to avoid running the application until termination
 - ▶ Provide ML-assisted hints to help developers find root causes of scalability bottlenecks

That's all folks!

<http://dca.ufrn.br/~samuel>

